Text classification into the large IPTC Media Topics taxonomy is challenging when relying solely on dense document embeddings, which often miss explicit factual context for rare named entities. This thesis investigates enhancing a multilingual, multi-label news classifier by injecting real-world entity knowledge. We propose a modular late-fusion architecture that combines standard text embeddings with structured entity representations derived from Wikidata and Wikipedia. We systematically evaluate multiple embedding sources, feature fusion techniques, and pooling strategies. The results confirm that entity-aware models consistently outperform the text-only baseline. The best configuration, which uses a learned attention mechanism to aggregate multilingual entity vectors, achieved a 1.7% improvement in Micro F_1 and a 2.4% increase in Macro F_1. This improvement was particularly substantial in specific domains, such as sports, showing that structured entity knowledge effectively addresses data sparsity and improves practical classification performance.